


# The normal distribution

INTRODUCTION TO STATISTICS IN R




**Maggie Matsui**  
Content Developer, DataCamp


# What is the normal distribution?




# Symmetrical



# Area = 1




# Curve never hits 0



# Described by mean and standard deviation

Mean: 20


Standard deviation: 3



*Standard normal distribution*

Mean: 0


Standard deviation: 1



# Described by mean and standard deviation

Mean: 20

Standard deviation: 3



*Standard normal distribution*

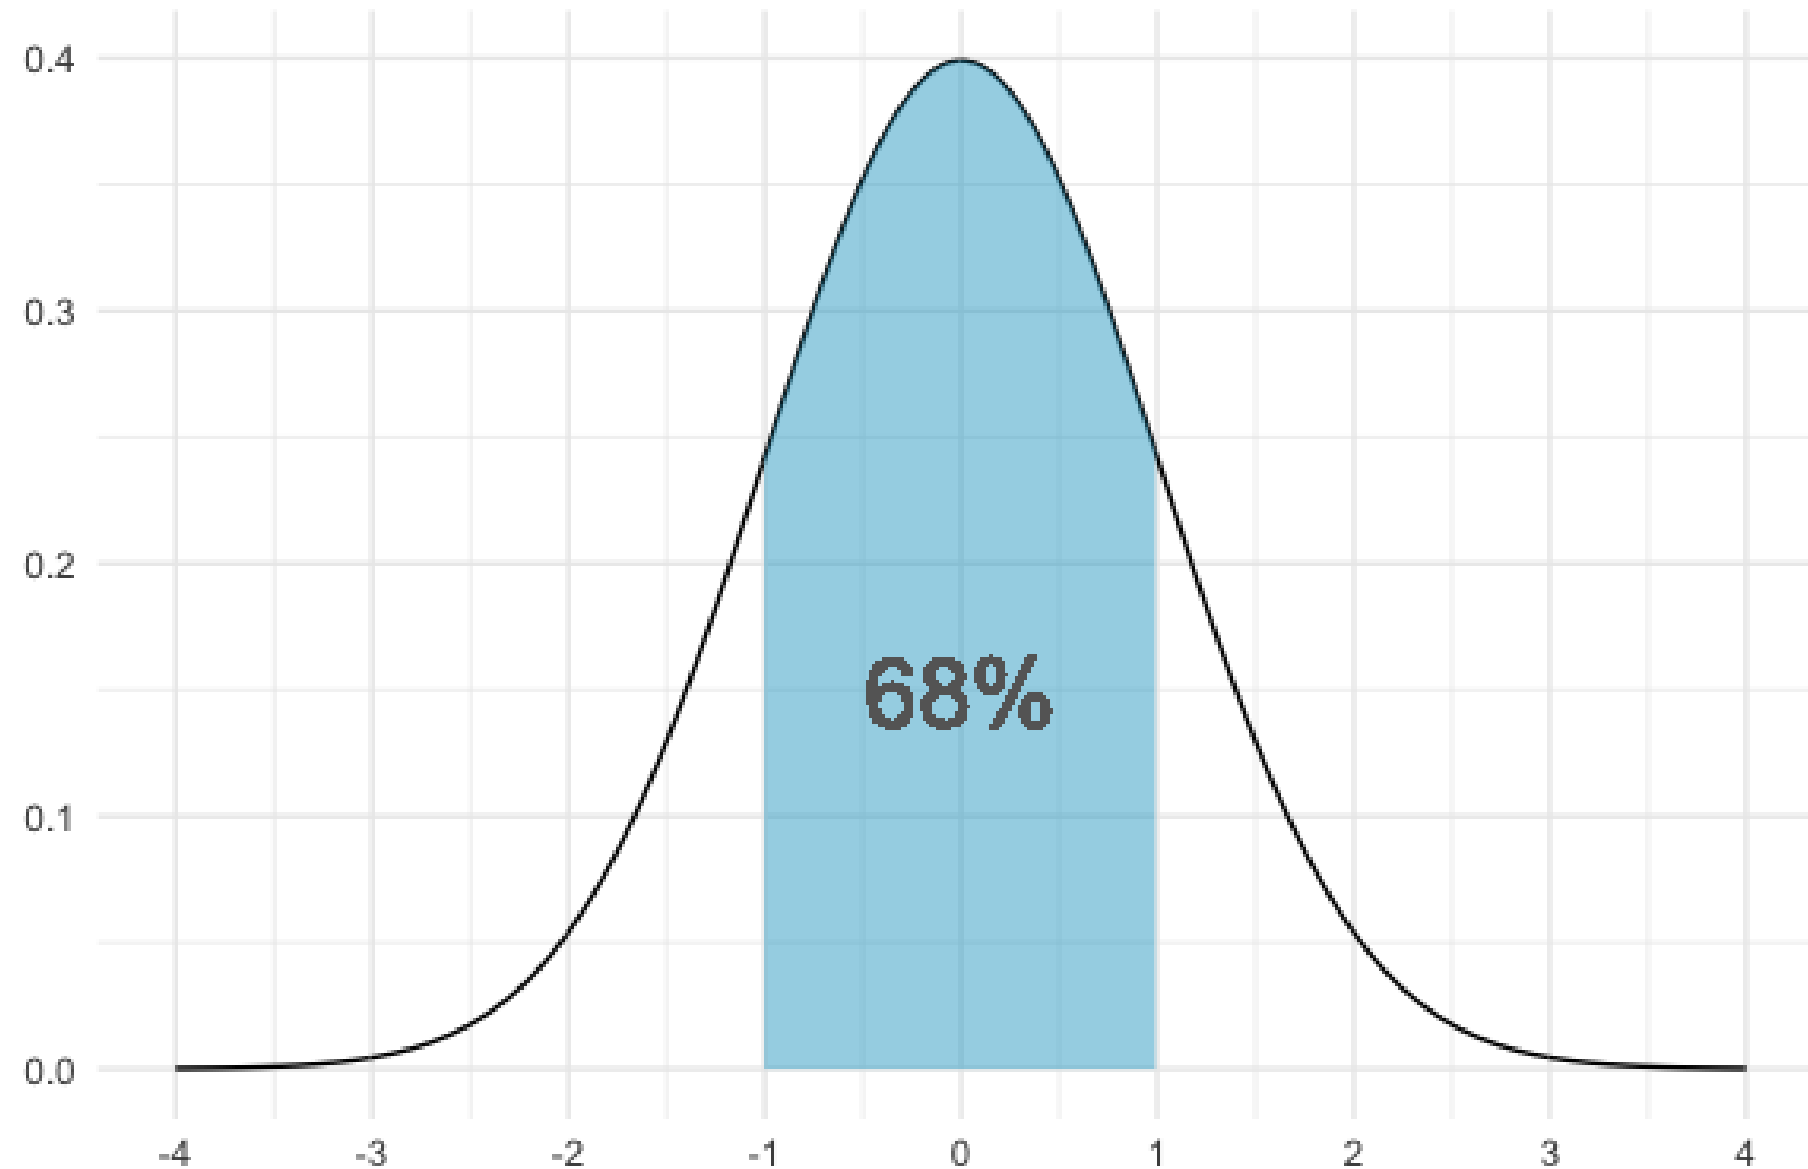
Mean: 0

Standard deviation: 1



# Areas under the normal distribution


68% falls within 1 standard deviation






# Areas under the normal distribution

95% falls within 2 standard deviations




# Areas under the normal distribution

99.7% falls within 3 standard deviations




# Lots of histograms look normal

## Normal distribution




## Women's heights from NHANES



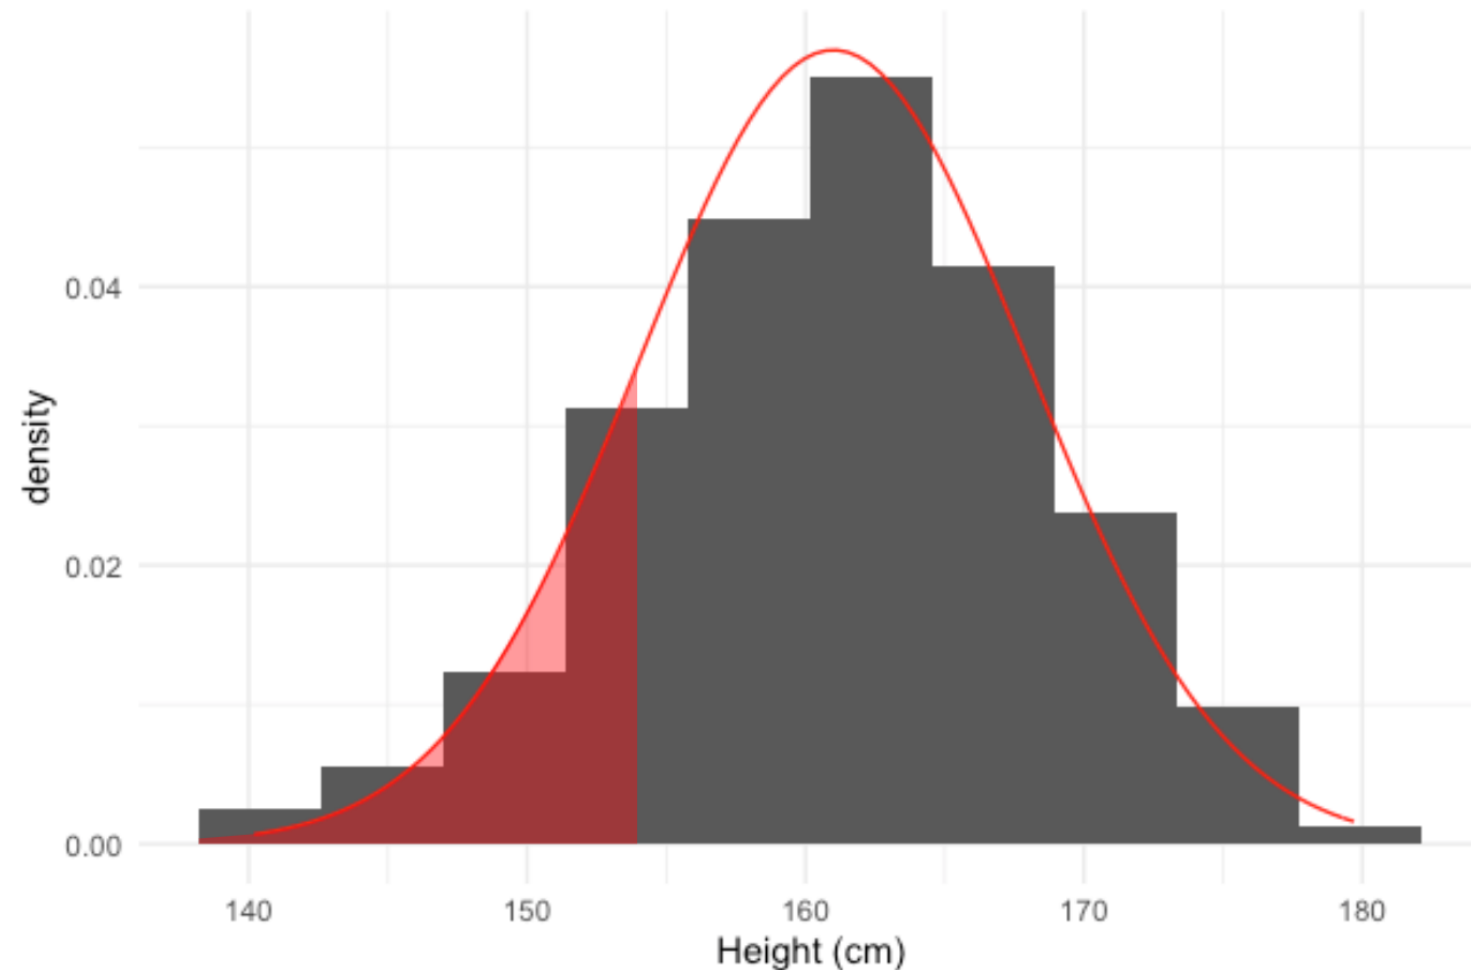
*Mean: 161 cm*

*Standard deviation: 7 cm*

# Approximating data with the normal distribution



# What percent of women are shorter than 154 cm?




```
pnorm(154, mean = 161, sd = 7)
```

```
0.159
```

*16% of women in the survey are shorter than 154 cm*


# What percent of women are taller than 154 cm?



```
pnorm(154, mean = 161, sd = 7,  
      lower.tail = FALSE)
```


```
0.8413447
```

# What percent of women are 154-157 cm?



```
pnorm(157, mean = 161, sd = 7) - pnorm(154, mean = 161, sd = 7)
```

# What percent of women are 154-157 cm?

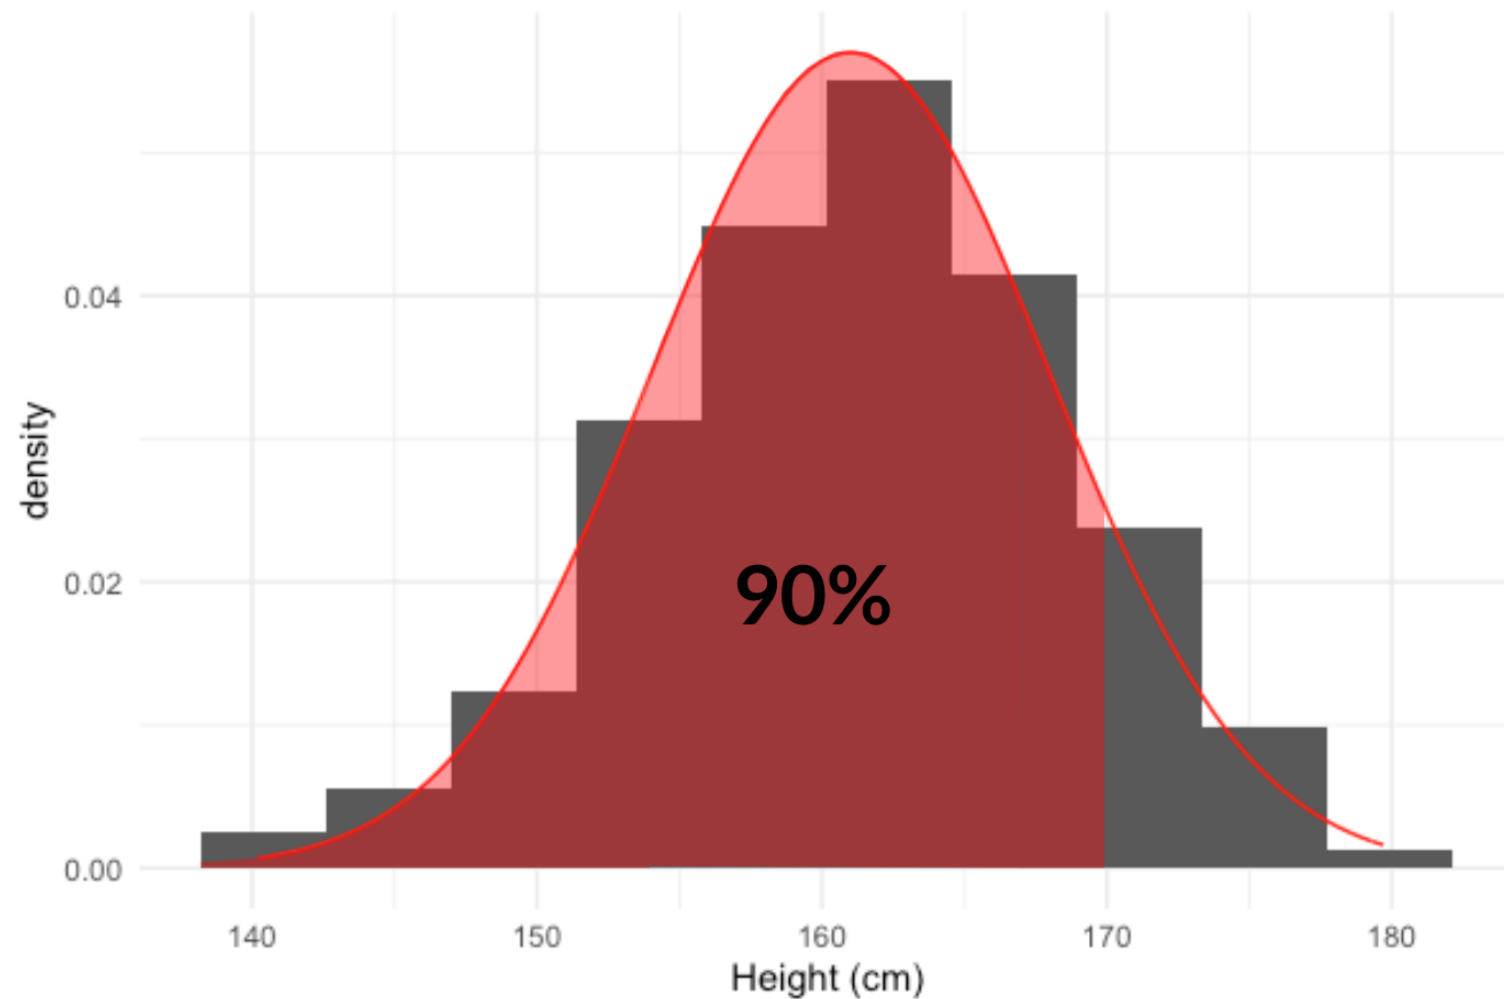


```
pnorm(157, mean = 161, sd = 7) - pnorm(154, mean = 161, sd = 7)
```

```
0.1252
```




# What height are 90% of women shorter than?



```
qnorm(0.9, mean = 161, sd = 7)
```

```
169.9709
```

# What height are 90% of women taller than?



```
qnorm(0.9,  
      mean = 161,  
      sd = 7,  
      lower.tail = FALSE)
```

152.03

# Generating random numbers

```
# Generate 10 random heights  
rnorm(10, mean = 161, sd = 7)
```

```
159.35 157.34 149.85 156.75 163.53 156.33 157.22 171.44 158.10 170.12
```

# Let's practice!

INTRODUCTION TO STATISTICS IN R

# The central limit theorem

INTRODUCTION TO STATISTICS IN R



**Maggie Matsui**  
Content Developer, DataCamp

# Rolling the dice 5 times

```
die <- c(1, 2, 3, 4, 5, 6)
# Roll 5 times
sample_of_5 <- sample(die, 5,
                      replace = TRUE)
sample_of_5
```

```
1 3 4 1 1
```

```
mean(sample_of_5)
```

```
2.0
```



# Rolling the dice 5 times

```
# Roll 5 times and take mean  
sample(die, 5, replace = TRUE) %>% mean()
```

```
4.4
```

```
sample(die, 5, replace = TRUE) %>% mean()
```

```
3.8
```

# Rolling the dice 5 times 10 times

*Repeat 10 times:*

- Roll 5 times
- Take the mean


```
sample_means <- replicate(10, sample(die, 5, replace = TRUE) %>% mean())  
sample_means
```

```
3.8 4.0 3.8 3.6 3.2 4.8 2.6 3.0 2.6 2.0
```



# Sampling distributions


*Sampling distribution of the sample mean*



# 100 sample means


```
replicate(100, sample(die, 5, replace = TRUE) %>% mean())
```

```
2.8 3.2 1.8 4.6 4.0 2.8 4.4 2.4 3.4 2.8 4.2 3.4 ... 2.2 3.8 3.6 3.8 4.4 4.8 2.4
```




# 1000 sample means

```
sample_means <- replicate(1000, sample(die, 5, replace = TRUE) %>% mean())
```



# Central limit theorem


The sampling distribution of a statistic becomes closer to the normal distribution as the number of trials increases.



\* *Samples should be random and independent*

# Standard deviation and the CLT

```
replicate(1000, sample(die, 5, replace = TRUE) %>% sd())
```



# Proportions and the CLT


```
sales_team <- c("Amir", "Brian", "Claire", "Damian")  
sample(sales_team, 10, replace = TRUE)
```

```
"Claire" "Brian" "Brian" "Brian" "Damian" "Damian" "Brian" "Brian"  
"Amir"   "Amir"
```

```
sample(sales_team, 10, replace = TRUE)
```

```
"Amir"   "Amir"   "Claire" "Amir"   "Amir"   "Brian"  "Amir"   "Claire"  
"Claire" "Claire"
```

# Sampling distribution of proportion



# Mean of sampling distribution


```
# Estimate expected value of die  
mean(sample_means)
```

```
3.48
```

```
# Estimate proportion of "Claire"s  
mean(sample_props)
```

```
0.26
```

- Estimate characteristics of unknown underlying distribution



- More easily estimate characteristics of large populations



# Let's practice!

INTRODUCTION TO STATISTICS IN R

# The Poisson distribution

INTRODUCTION TO STATISTICS IN R



**Maggie Matsui**  
Content Developer, DataCamp

# Poisson processes

- Events appear to happen at a certain rate, but completely at random
- Examples
  - Number of animals adopted from an animal shelter per week
  - Number of people arriving at a restaurant per hour
  - Number of earthquakes in California per year




# Poisson distribution

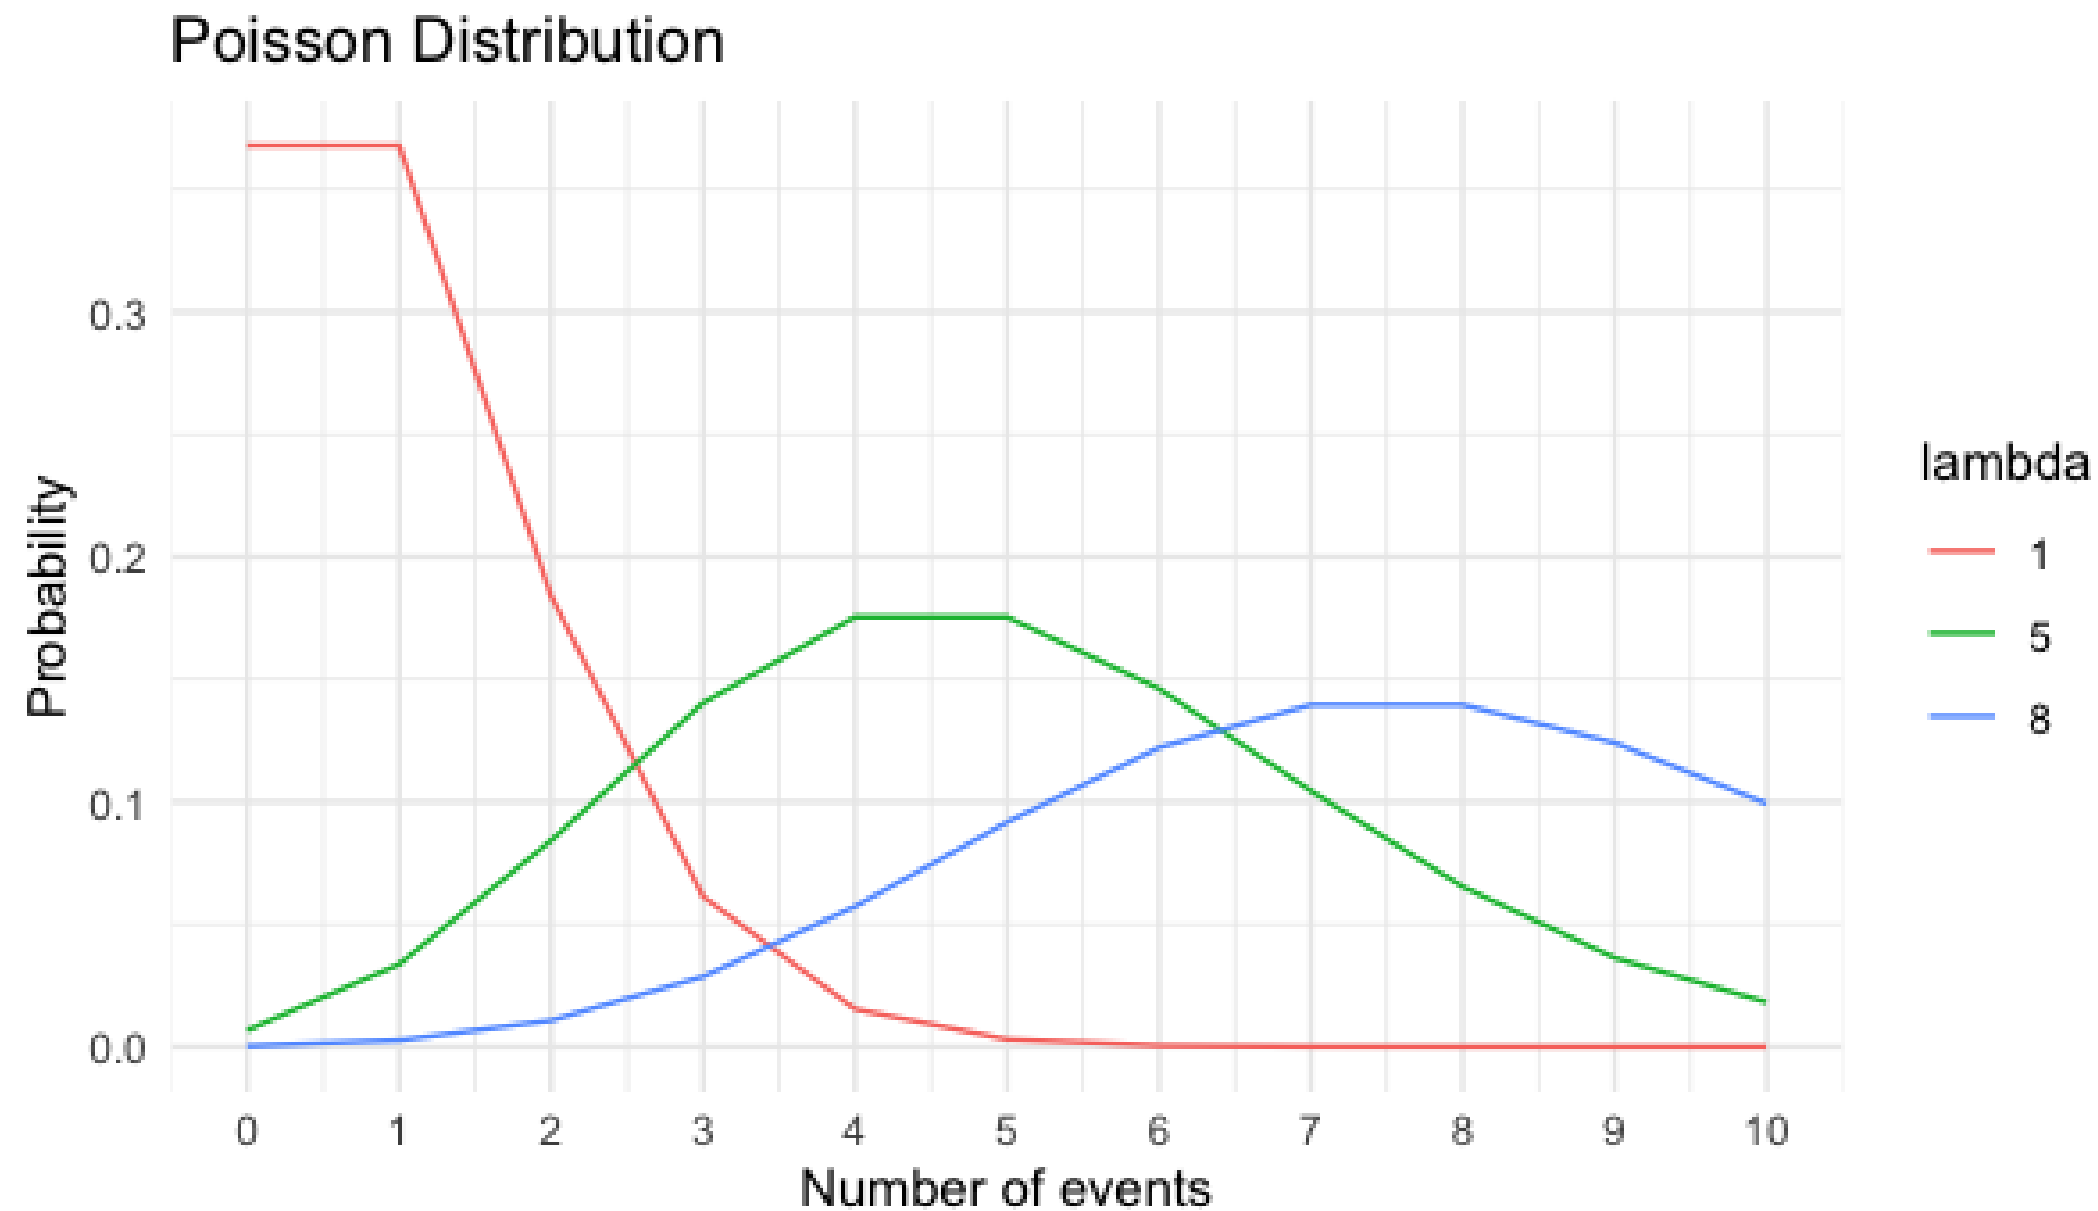
- Probability of some # of events occurring over a fixed period of time
- Examples
  - Probability of  $\geq 5$  animals adopted from an animal shelter per week
  - Probability of 12 people arriving at a restaurant per hour
  - Probability of  $< 20$  earthquakes in California per year

# Lambda ( $\lambda$ )

- $\lambda$  = average number of events per time interval
  - Average number of adoptions per week = 8



# Lambda is the distribution's peak



# Probability of a single value

If the average number of adoptions per week is 8, what is  $P(\# \text{ adoptions in a week} = 5)$ ?

```
dpois(5, lambda = 8)
```

```
0.09160366
```

# Probability of less than or equal to

If the average number of adoptions per week is 8, what is  $P(\# \text{ adoptions in a week} \leq 5)$ ?

```
ppois(5, lambda = 8)
```

```
0.1912361
```



# Probability of greater than

If the average number of adoptions per week is 8, what is  $P(\# \text{ adoptions in a week} > 5)$ ?

```
ppois(5, lambda = 8, lower.tail = FALSE)
```

```
0.8087639
```

If the average number of adoptions per week is 10, what is  $P(\# \text{ adoptions in a week} > 5)$ ?

```
ppois(5, lambda = 10, lower.tail = FALSE)
```

```
0.932914
```


# Sampling from a Poisson distribution

```
rpois(10, lambda = 8)
```

```
13  6 11  7 10  8  7  3  7  6
```

# The CLT still applies!

Distribution of sample means from Poisson distribution ( $\lambda = 8$ )



# Let's practice!

INTRODUCTION TO STATISTICS IN R

# More probability distributions

INTRODUCTION TO STATISTICS IN R



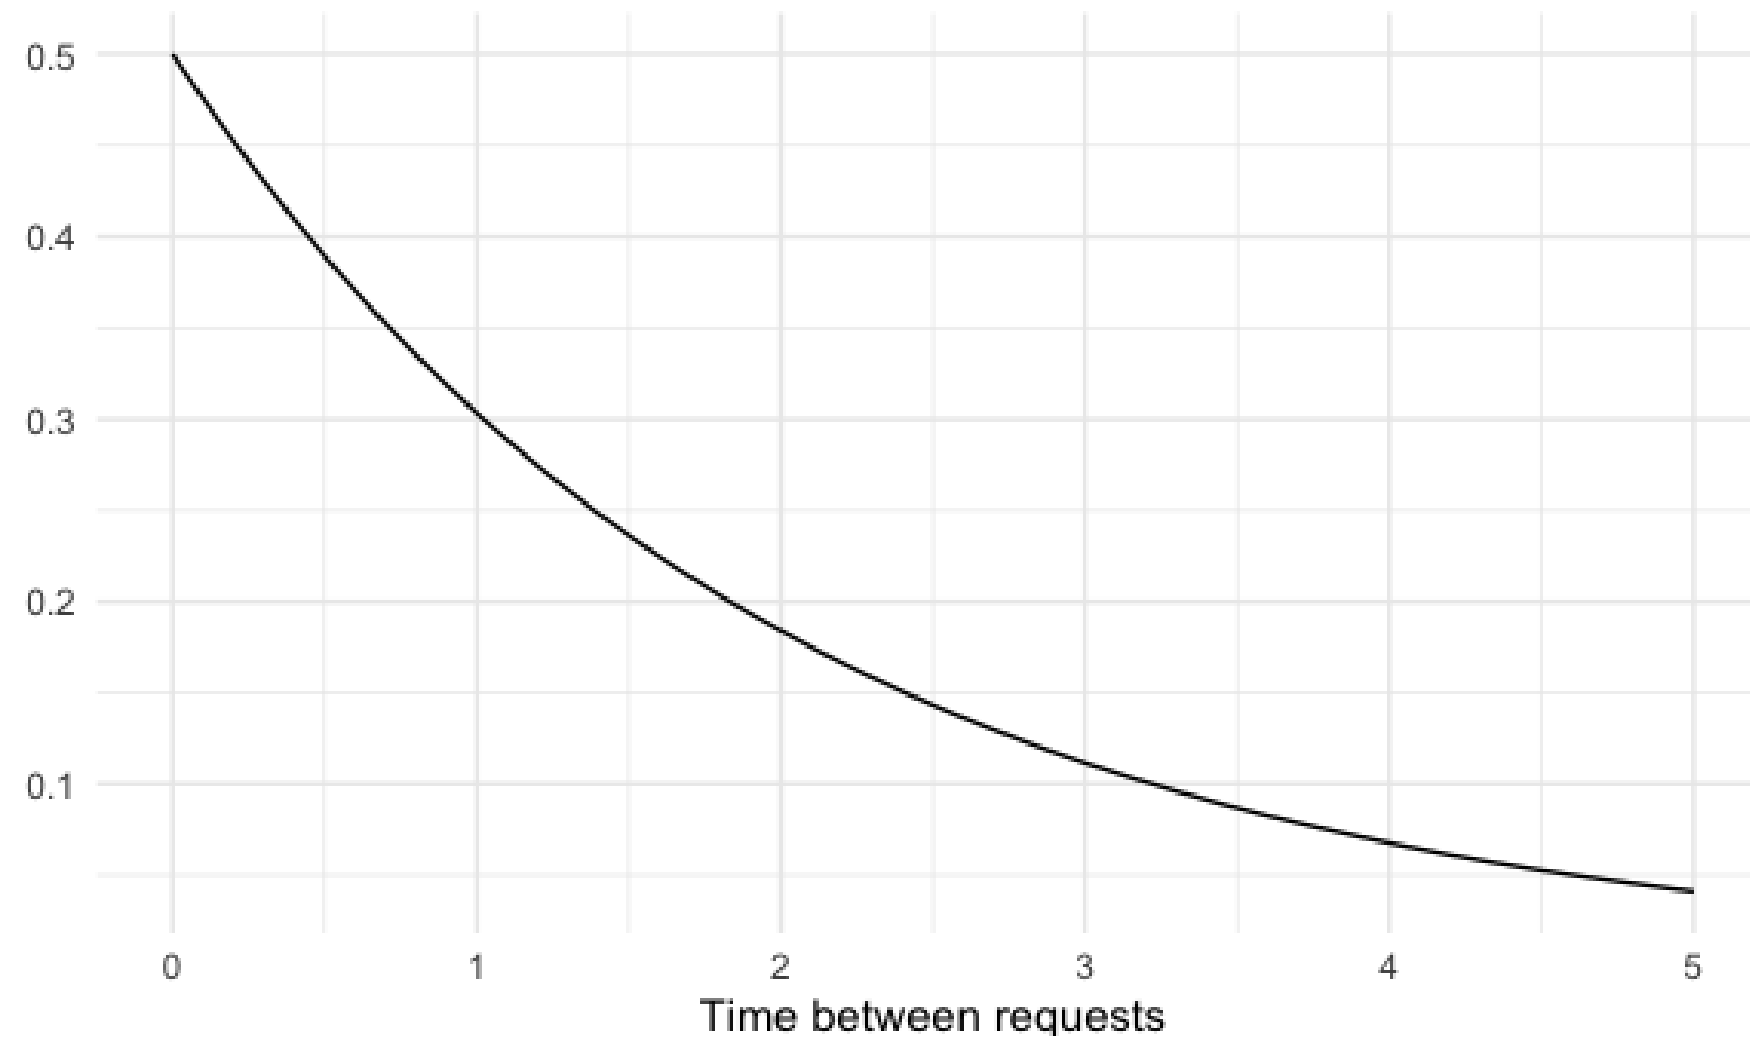
**Maggie Matsui**  
Content Developer, DataCamp

# Exponential distribution

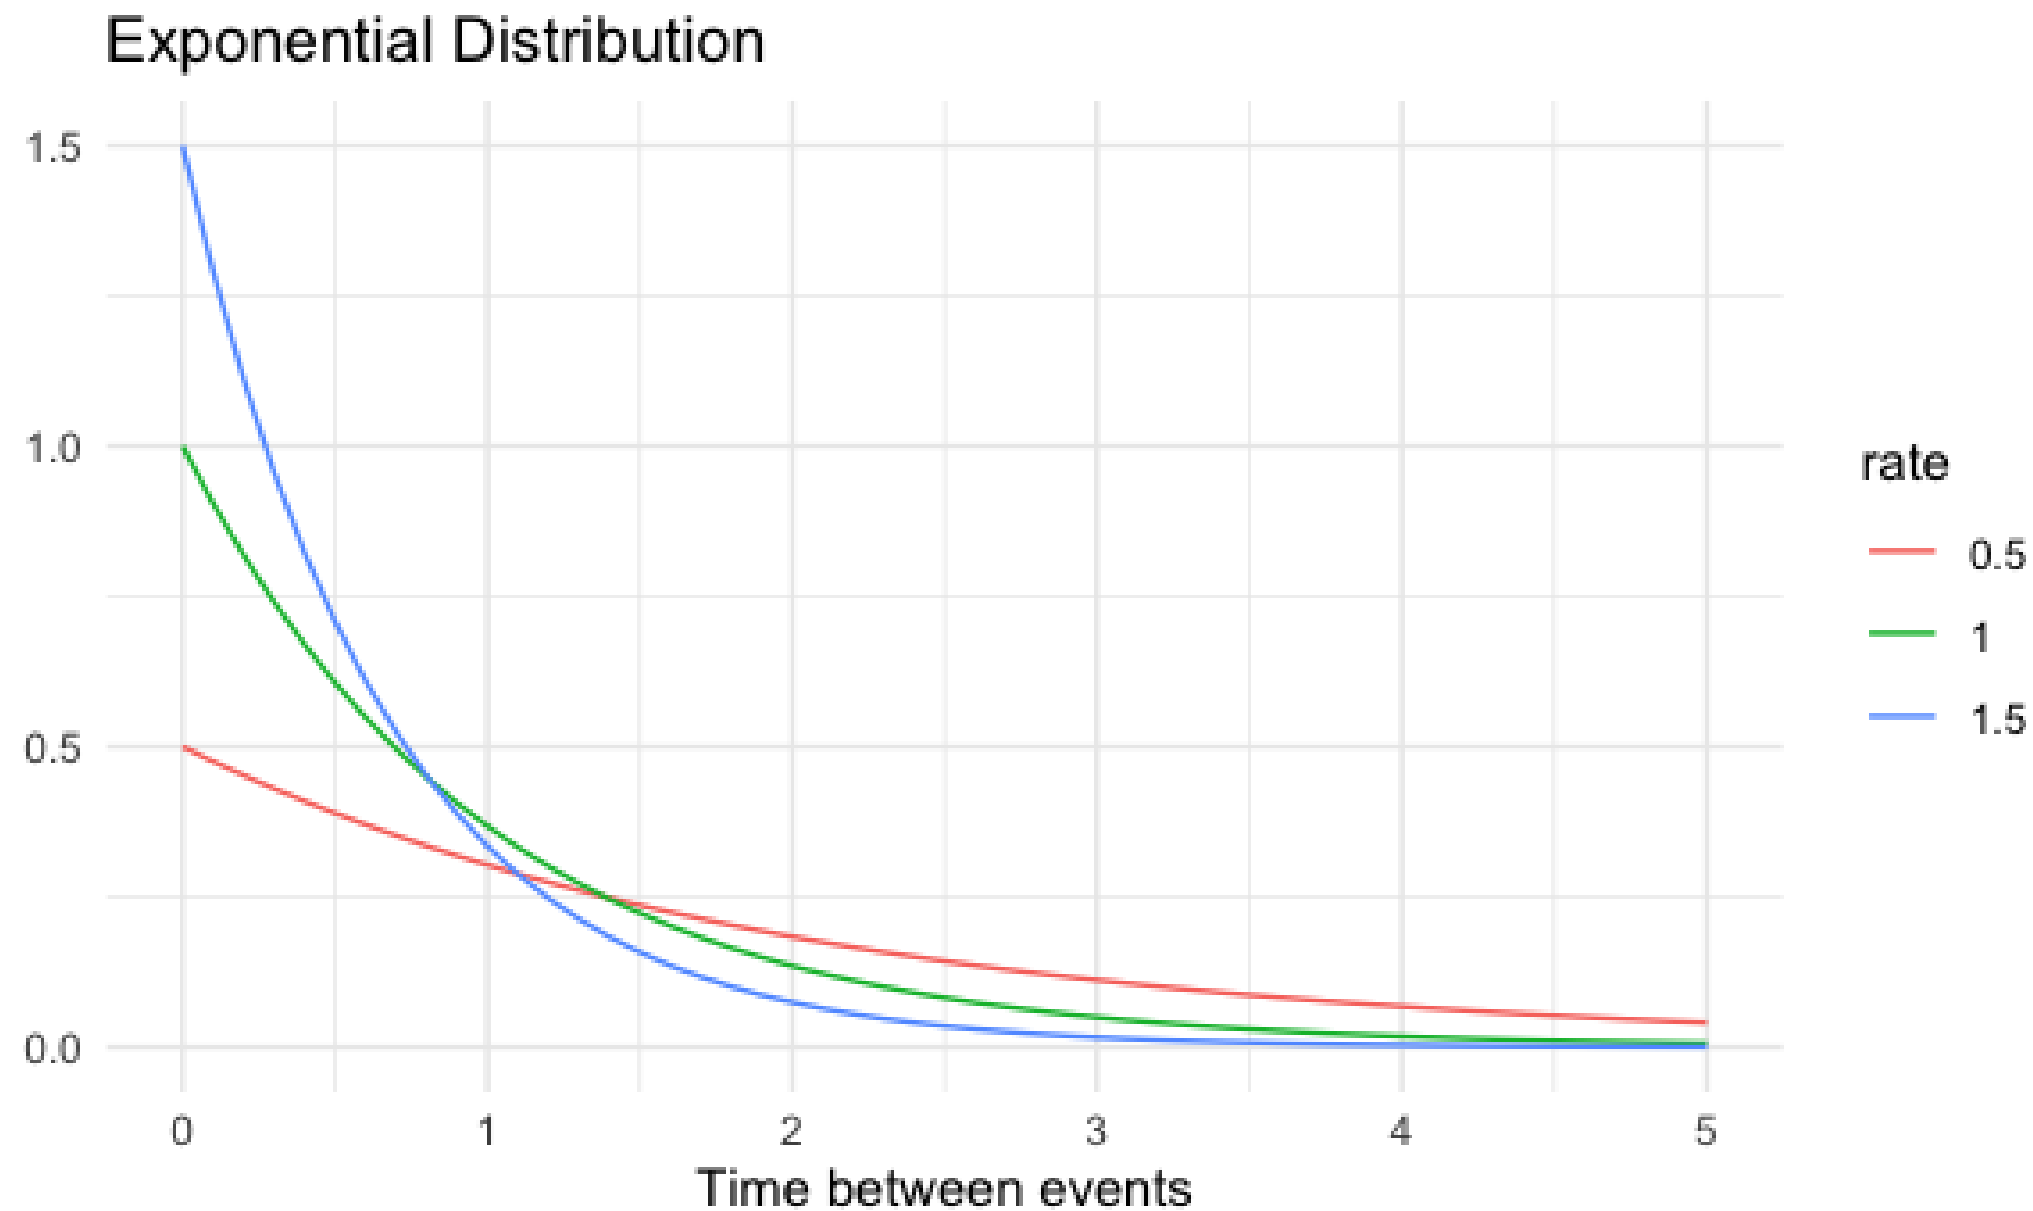
- Probability of time between Poisson events
- Examples
  - Probability of  $> 1$  day between adoptions
  - Probability of  $< 10$  minutes between restaurant arrivals
  - Probability of 6-8 months between earthquakes
- Also uses lambda (rate)
- Continuous (time)

# Customer service requests

- On average, one customer service ticket is created every 2 minutes
  - $\lambda = 0.5$  customer service tickets created each minute



# Lambda in exponential distribution





# How long until a new request is created?

$$P(\text{wait} < 1 \text{ min}) =$$

```
pexp(1, rate = 2)
```

```
0.8646647
```

$$P(\text{wait} > 3 \text{ min}) =$$

```
pexp(3, rate = 2, lower.tail = FALSE)
```

```
0.002478752
```

$$P(1 \text{ min} < \text{wait} < 3 \text{ min}) =$$

```
pexp(3, rate = 2) - pexp(1, rate = 2)
```

```
0.1328565
```

# Expected value of exponential distribution

In terms of rate (Poisson):


- $\lambda = 0.5$  requests per minute

In terms of time (exponential):

- $1/\lambda = 1$  request per 2 minutes


# (Student's) t-distribution

- Similar shape as the normal distribution




# Degrees of freedom

- Has parameter degrees of freedom (df) which affects the thickness of the tails
  - Lower df = thicker tails, higher standard deviation
  - Higher df = closer to normal distribution



# Log-normal distribution

- Variable whose logarithm is normally distributed
- Examples:
  - Length of chess games
  - Adult blood pressure
  - Number of hospitalizations in the 2003 SARS outbreak



# Let's practice!

INTRODUCTION TO STATISTICS IN R